

Poker Zero Lightning Presentation

Problem Statement and Motivation

Objective

Develop a reasoning model optimized for **no-limit hold'em poker**. Success in poker extends beyond monetary gain—it signifies advancements in:

- Reasoning under incomplete information
- Adversarial decision-making
- Strategic adaptation

Applications

Poker AI has significant implications in:

- **Game Theory**
- **Economic Modeling**
- **Negotiation**

Prior Work

Several notable AI-driven poker models include:

- **Poker Bench** – Trained LLMs to become professional poker players
- **PokerGPT** – Lightweight solver leveraging a large language model
- **Pluribus** – Demonstrated multiplayer, near-GTO strategies

Data

Instruction	Output
<p>You are a specialist in playing 6-handed No Limit Texas Hold'em.</p> <p>The following will be a game scenario, and you need to make the optimal decision.</p>	Bet 18
<p>Game Summary:</p>	
<p>- Small Blind: 0.5 chips</p>	
<p>- Big Blind: 1 chip</p>	
<p>- Everyone started with 100 chips</p>	
<p>- Player positions: UTG, HJ, CO, BTN, SB, BB</p>	
<p>- Your Position: HJ</p>	
<p>- Your Hand: [King of Diamond, Jack of Spade]</p>	

Technical Approach

Unsloth

Unsloth is an open-source Python framework that speeds up the process of fine-tuning and using LLMs.

- Optimized PyTorch code
- Handwriting GPU kernels to speed up inference
- Better memory utilization via typecasting
- Allows us to fine-tune 8B parameter models on just a Colab T4 GPU

GRPO

GRPO (Group Relative Policy Optimization) is a reinforcement learning algorithm developed by DeepSeek.

- Trains a model to optimize a reward function instead of training a model solely on next-token prediction (which simply teaches it to mimic data)
- Uses group-based comparisons to improve performance instead of after every trial
- Calculates advantage and updates policy to increase likelihood of better actions
- Uses KL Divergence constraint to prevent drastic changes in policy
- Overall objective maximizes cumulative reward with stable policy updates

LoRA

Low Rank Adaptation (LoRA) is a method for fine-tuning large models efficiently.

- It works by introducing low-rank matrices into pretrained model layers.
- This approach allows significant performance gains with a minimal number of additional parameters.

Mathematical Formulation of LoRA

- **Pretrained Weight Matrix:** Let $W \in \mathbb{R}^{d \times k}$ be a pretrained weight matrix.
- **Low-Rank Decomposition:** LoRA approximates the weight update ΔW as:
$$\Delta W = BA$$

where:
 - $B \in \mathbb{R}^{d \times r}$
 - $A \in \mathbb{R}^{r \times k}$
 - $r \ll \min(d, k)$
- **Adapted Weights:** The updated weight matrix becomes: $W' = W + BA$

Initial Results

- Successfully fine tuned our model in Colab, showing Unsloth's capabilities in speeding up fine-tuning
- **PyPokerEngine:** A library for simulating poker games with AI bots
 - Used PyPokerEngine to make our model play against itself
- Initial testing through sanity check on output

Next Steps

PyPokerEngine

- **More Players** Poker environment is only heads up (2 players). Our goal is to play 6 handed, which would have more interesting applications because heads-up Texas Hold'em is a solved game

Additional Training

- **Self-Play** Model competes against previous versions, gradual improvement and adaptation to exploitative strategies
- **Training Pipeline** Feed data back into our model for reinforcement learning training. Eliminates need for a premade dataset and allows the model to learn optimal strategies by itself

Thank You!

Questions?