

Poker Zero Final Presentation

Problem Definition

- **What are we solving?**
 - We aim to build a **reasoning model** for No-Limit Hold'em poker capable of decision-making under **incomplete information** and **adversarial conditions**.
- **Importance:**
 - Real-time strategic thinking, risk assessment, and adaptation.
 - Decision-making under uncertainty
- **Success Criteria**
 - **Win rate** (hands won, stack size, profit over time)
 - **Performance against GTO (Game Theory Optimal) strategies**

Prior Work

- Poker solvers often play game theory optimal poker, which is limited
 - Nash equilibria is hard to compute for multi-way zero sum games
 - Only able to calculate on a limited set of scenarios, since poker has a very large game tree
 - Won't take advantage of imperfect opponent

Using ML and LLMs

- State of the art poker bot: **Pluribus**
 - Uses self-play to iteratively converge close to the Nash Equilibrium
 - Leverages Monte Carlo Counterfactual Regret Minimization
- Transformer models like ChatGPT / GPT-4 also don't play GTO
- LLMs use less compute/resource consumption than CFR
- Can receive more information in the game tree

Modeling the Problem

- **Stochastic Nature of Poker:**

- Poker is inherently random, with incomplete information and unpredictable outcomes

$$EV = \sum_{h \in H} P(h) \times R(h)$$

- $P(h)$ is the probability of a hand outcome $R(h)$ is the corresponding reward.

- **Goal:** maximize total expected winnings over all rounds r

- **Actions:** Possible moves (bet, raise, fold, call)

Modeling the Problem

- Reinforcement Learning (RL) naturally suited for sequential decision making problems and long term problems
- Setting up a basic framework
 - **States:** Game configurations, including hole cards, community cards, betting history, stack sizes
 - **Actions:** Possible moves (bet, raise, fold, call) & associated amounts
 - **Rewards:** Based on how good the actions are

Optimization Methods

- **Policy Gradient Methods:**

Directly optimize the policy by maximizing the expected reward.

- **PPO (Proximal Policy Optimization):**

Balances exploration and exploitation with clipping or penalty methods to ensure stable updates.

- **GRPO (Group Relative Policy Optimization):**
 - **Multi-Agent Environment**
Adjusts an agent's policy relative to a group baseline or relative to other agents' policies.
 - **Reward-Centric Updates:**
Focuses on adjusting policies based on long-term reward estimates.
 - **Empirical Advantages:**
Demonstrates superior performance compared to PPO and TRPO in several benchmarks.
- **Relevance to Poker:**
Enables strategic adaptation and robust performance in complex, adversarial settings.

Unslloth

- **Purpose:**

An open-source Python framework optimized for fast fine-tuning and deployment of large language models.

- **Key Features:**

- Optimized PyTorch code
- Handwritten GPU kernels to speed up inference
- Significantly reduced memory usage
- Allows us to fine-tune 8B parameter models on a Colab T4 GPU

LoRA (Low Rank Adaptation)

- **Concept:**

Introduces low-rank matrices into pretrained model layers to achieve efficient fine-tuning.

- **Benefit:**

Significant performance gains with a minimal increase in parameters—ideal for adapting large models in resource-constrained environments.

Mathematical Formulation of LoRA

- **Pretrained Weight Matrix:**

Let $W \in \mathbb{R}^{d \times k}$ be a pretrained weight matrix.

- **Low-Rank Update:**

Approximate the weight update as:

$$\Delta W = BA$$

where:

- $B \in \mathbb{R}^{d \times r}, A \in \mathbb{R}^{r \times k}, r \ll \min(d, k)$

- **Adapted Weights:**

The new weight matrix is given by:

$$W' = W + BA$$

Weight Update, Factorization, and Training Efficiency

- **Factorization & Weight Update:**
 - Originally: $O(dk)$ parameters
 - With LoRA: $O(r(d + k))$ parameters
- **Training via a New Loss Function:**
 - $f(B, A) = L(W_0 + BA)$
 - Optimize this loss by taking gradient steps with respect to B and A
 - W_0 is frozen
- **In our implementation**
 - r is a hyperparameter
 - We choose $r = 32$, about $\frac{59,867,136}{3,000,000,000}$ ($\sim 2\%$) of parameters are trainable

Optimization & Reward Functions

- **Reward Function Design:**
 - **For Initial Training:**
 - **Negative Reward:** Apply penalties for outputs that violate constraints.
 - **Zero Reward:** No reward for clearly incorrect moves.
 - **Partial Rewards:**
 - Reward for executing a correct action.
 - Additional reward for an almost correct action (e.g., bet size within $\pm 20\%$ of the optimal).
 - **Maximum Reward:**
 - Full reward for both the correct action and optimal bet sizing.

- **Why This Approach?**

- Allows gradual, nuanced learning instead of an all-or-nothing reward.
- Helps the model learn the subtleties of decision-making in an environment where perfect play is rare.

Selecting GRPO & Refining the Model

- **Algorithm Choice: Group Relative Policy Optimization (GRPO)**
 - **Justification:**
 - Suitable for poker's continuous and complex environment where isolated wins do not ensure overall success.

- **Tuning Procedure & Hyperparameters:**
 - **Reward Function Tuning:**
 - Began with rewards only for exact matches, but feedback was sparse.
 - Introduced partial credit for near-miss outputs (e.g., valid poker moves, near-optimal bet amounts).
 - **Hyperparameter Exploration:**
 - Systematic grid search over reward thresholds and learning rates.
 - Iterative refinement based on model performance and stability.



- **Implementation Choices:**



- Leveraged the **Unslloth** framework to optimize GRPO training.
- Utilized PyTorch for model development and integration, taking advantage of its efficient computation and GPU support.

Self Play


- **PyPokerEngine:** A library for simulating poker games with AI bots
 - Used PyPokerEngine to make our model play against itself
- Model plays against earlier iterations of itself
 - Uses the output to generate additional training data
- Challenges: Took a long time to run
 - 6 instances of our model have to conduct inference and give outputs

Results






Rewards for Self Play Round 1



Rewards for Self Play Round 5



Results Overall

- **Profit Rate:** Consistent improvement in win rates against older iterations of the models

How do your results compare to baseline methods or the literature?

- **Performance Metrics:**
 - **Reward Rate:** Upward trend of rewards over time

Demo

Showcase a demo or compelling visualization if applicable.

Compare expected progress with actual progress. Explain discrepancies.

Project Reflection

Limits Encountered & Adaptations

- **Computational Resources:**
 - Initially limited to a T4 GPU on Colab, leading to frequent disconnections and slow iteration.
 - Challenges in accessing scalable GPU resources on platforms like Google Cloud.
- **Impact on Model Training:**
 - Slow training and iteration speeds forced us to adjust our training framework.
 - Required tuning reward functions to provide a denser, more continuous feedback signal.

- **Adaptation Strategies:**

- **Unslloth** played a critical role in speeding up our training cycles.
- Optimization of training loops and hyperparameter searches to work within computational constraints.
- Adoption of incremental learning strategies to mitigate resource limitations while still aiming for optimal performance.

Technical/Conceptual difficulties

- Understanding the complexity of poker strategies and how to model them effectively.
- Implementing reinforcement learning algorithms, especially in the context of self-play.

What part of the project workflow was easier than expected? Harder?

Easier

- Implementing basic reinforcement learning algorithms using **unsloth** directly within Colab.

Harder

- Debugging and tuning reinforcement learning models to converge effectively in self-play scenarios.

How the project evolved

- Trained initial model using **GRPO** and attempted self-play reinforcement learning with **PPO** to generate initial neural layers.
- **Challenge:** Model wasn't converging.
 - Shifted to using the self-play environment to generate additional training data for **GRPO**.
- Focused more on **self-play**, with multiple iterations of PokerZero playing each other to measure performance improvements.

For The Future

- Rewards come from correct moves AND winning the pot, so model receives less rewards for folds, even if they are correct
- Add additional analysis and provide rewards for good folds
- Playing against other models besides our own to evaluate performance
- Need way more training time and way more compute to achieve our goal

How did AI tools assist your project?

Specific Examples:

- **Literature Review:** Helped in understanding initial concepts and strategies for reinforcement learning.
- **Debugging:** Assisted in understanding complex algorithms and generating code