

Implementation Updates

Training Infrastructure Improvements

Unslloth Optimization for T4 GPUs

- Successfully migrated training from A100 GPUs to T4 GPUs using Unslloth's notebook optimizations
- Enables more cost-effective and accessible training
- Maintains model quality with extended training times

Response Structure Revisions

- Implemented a structured output format for LLM
- Encourages explicit reasoning rather than only providing decisions
- Reward Functions now evaluate:
 - **Structure Compliance:** Rewards proper tag structure
 - **Reasoning Quality:** Assesses depth, relevance, and logic
 - **Decision Alignment:** Ensures reasoning supports the final decision
- More lenient rewards for faster convergence

Reinforcement Learning Implementation

Self-Play with PPO

- Implemented **Proximal Policy Optimization (PPO)** for self-play using PyPokerEngine
- Features:
 - 6-player poker games with model instances competing against each other
 - Gameplay data collection for continuous improvement
 - Competitive reinforcement learning for decision refinement

GRPO Training Challenges

- Integrated **Group Relative Policy Optimization (GRPO)** for enhanced learning
- Challenges faced:
 - **Computational Intensity:** Significant time required for updates
 - **Complexity Management:** Poker's vast state space adds difficulty
 - **Long Runtimes:** Full training cycles are resource-intensive
- Further optimization needed for practical training cycles

Enhanced Reward Functions

Verbal Descriptions

- **Structure Compliance:**

Models are rewarded for adhering to predefined tag formats. This ensures consistent output, making evaluation and debugging easier.

- **Reasoning Quality:**

Reasoning sections are assessed for relevance and logical consistency. More detailed and insightful reasoning results in higher rewards.

- **Decision Alignment:**

Models earn rewards when the final decision logically follows from the reasoning. Misalignment is penalized to reinforce accurate decision-making.

- **Granular Feedback:**

Rewards are applied incrementally, reducing punishment for minor errors while

Next Steps

Training Efficiency

- Explore:
 - Smaller models for faster fine-tuning
 - Algorithmic improvements to reduce compute intensity

Reward Function Refinements

- Detect and mitigate reward hacking
- Enhance feedback for better reasoning quality

Model Evaluation

- Current Method: Track average chip differences in games
- Future Exploration:

